


Whole-model BF16 training vs LoRA · save mode · RTX 5090



4 layers · width 768 · 2048 tokens/step · frozen base/head for LoRA · FP32 factors, gradients and AdamW moments